


AI Agent Layer on a Risk-Decisioning System

(anonymized reference architecture)



Design principles: ① Additive, not invasive — existing models & policy engine stay untouched
② Grounded reasoning — agents call tools over real system outputs, never free-generate facts
③ AI adds judgment & language; the existing system remains the single source of truth
④ Full auditability — every agent decision and generated narrative is logged & traceable